
TODO

TODO

BACHELOR THESIS

by

TODO

in fulfillment of requirements for degree

BACHELOR OF SCIENCE (B.Sc.)

submitted to

RHEINISCHE FRIEDRICH-WILHELMS-UNIVERSITÄT BONN

INSTITUT FÜR INFORMATIK IV

ARBEITSGRUPPE FÜR IT-SICHERHEIT

in degree course

COMPUTER SCIENCE (M.Sc.)

First Supervisor: **TODO**
University of Bonn

Second Supervisor: **TODO**
TODO

Sponsor: **TODO**
University of Bonn

TODO

ABSTRACT

TODO

CONTENTS

1	INTRODUCTION	1
2	BACKGROUND	3
2.1	CPU	3
2.2	Out-of-order execution	3
2.3	Speculative execution	5
2.4	Meltdown and Spectre	5
2.4.1	Meltdown	5
2.4.2	Spectre	6
2.4.3	Mitigations	6
3	SPECIFICATION OF OUR TASK	7
4	CPU EMULATOR/ BACKEND	8
4.1	CPU Components and our equivalents/ models (10-11 pages)	8
4.1.1	CPU	8
4.1.2	Instructions and Parser	9
4.1.3	Data representation	9
4.1.4	CPU frontend	9
4.1.5	Memory	11
4.1.6	Execution Engine	11
4.2	Out of Order Execution (2 pages)	12
4.2.1	Issuing instructions	12
4.2.2	Executing instructions	13
4.2.3	Memory hazards	13
4.2.4	Fence instruction	14
4.3	Exceptions and Rollbacks (2-3 pages)	14
4.4	ISA (2 pages)	14
4.4.1	Default Instruction Set	15
4.5	Config Files (1 page)	17
5	USER INTERFACE AND USAGE	18
6	EVALUATION	19
7	CONCLUSION	20
8	CATCHY LAB TITLE	21
8.1	Abstract	21

8.2	Introduction (1 page)	21
8.3	Brief Theoretical Background (2-3 pages)	21
8.4	maybe Specification of the task (1 page)	22
8.5	CPU emulator/ backend (17-18 pages)	23
8.6	Our Visualisation and Usage/ Frontend (10 pages)	25
8.7	Demonstration/ evaluation (7 pages)	26
8.8	Conclusion (1 Page)	27
REFERENCES		28
LIST OF FIGURES		29

1 INTRODUCTION

Contrary to older processors, modern Intel CPUs implement a number of optimization techniques that increase their efficiency. One of which is the concept of out-of-order execution, which takes advantage of the mutual independence of instructions that would normally be executed sequentially. A second optimization technique is called speculative execution and involves the prediction of which branches are taken. With either technique, the CPU might encounter cases where the current CPU state must be rolled back to a previous one. For out-of-order execution, this happens if an instruction raises an exception (illegal memory access, for example). For speculative execution, this happens if a branch is mispredicted. A rollback causes some instructions that are currently being executed (in-flight) to continue execution for a short amount of time.

Even though rollbacks are meant to make sure in-flight instructions do not cause any lasting side effects on the microarchitectural state of the CPU, it was discovered that they can change the contents of caches and other buffers. The disclosure of both Spectre and Meltdown in early 2018 introduced a whole family of vulnerabilities that take advantage of both out-of-order and speculative execution to leak secrets over the processor's caches. And while the performance losses introduced by software and hardware mitigations are measurable, neither vulnerability can be exploited freely on a fully patched system. As a result, however, the process of trying to exploit one of the vulnerabilities for the sake of learning how they work in detail can be challenging. Apart from the software, which may be obtained by installing an older version of an operating system that does not implement any mitigations, one must also make sure their CPU is effected by the vulnerabilities and has not yet received any relevant microcode updates from Intel. Often times, this means that a user's personal computer does not meet these requirements.

We design and implement a graphical CPU emulator that supports single-step out-of-order and speculative execution and is vulnerable to select variants of Spectre and Meltdown. While it is a simplification in comparison to real hardware, the emulator allows its users to gain a better understanding of how exactly the two vulnerabilities work and can be exploited. Furthermore, the user may experiment with (ineffective) mitigations or implement their own microprograms that are executed once rollbacks are completed. We supply example programs that can be run by the emulator and serve both as an entry point for the user as well as the basis of our evaluation.

Firstly, chapter 2 briefly gives an overview of relevant components of vulnerable Intel CPUs, presents the concepts of out-of-order and speculative execution in greater detail, and introduces both Meltdown and Spectre to which the resulting emulator is supposed to be vulnerable. Secondly, chapter 3 further describes the target audience of the emulator and which variants of the vulnerabilities have been chosen, while chapter 4 documents the implementation of the emulator by describing each main component and explaining how rollbacks are implemented. Additionally, it contains an overview of the set of ISA instructions available to the user. Furthermore, chapter 5 explores the

graphical user interface by defining the goals its design is supposed to accomplish and documenting design choices and important features. Chapter 6 provides a demonstration of the emulator, which includes example programs, and determines how effective the implemented mitigations are. Finally, chapter 7 summarizes the other chapters, briefly reflects on how valuable the emulator might be to our target audience, and gives ideas for future improvements.

2 BACKGROUND

This chapter briefly covers the theoretical background needed to use the presented emulator and continue its development. The reader is assumed to have an understanding of elementary CPU concepts, such as pipelining and caching.

2.1 CPU

A CPU consists of a frontend, an execution engine, and a memory subsystem. As per Intel’s Skylake architecture [Wik], the frontend fetches the instructions, maintains a queue of instructions that are to be executed, and decodes them into simpler microinstructions, which are then communicated to the execution engine. Additionally, it is responsible for the branch prediction (see sec. 2.3). [Gru20]

The execution engine consists of multiple sets of execution units, each set being responsible for a specific type of microinstruction, such as loads, stores, or arithmetic. Further, the scheduler allows the execution units to work on independent instructions in parallel, while the reorder buffer makes sure that instructions retire in the correct order. A common data bus (CDB) connects the reorder buffer, scheduler, and execution units. Its purpose is further described in sec. 2.2. [Gru20]

Lastly, the memory subsystem handles all memory accesses of the execution units by maintaining caches and ensuring data is fetched from lower level caches or DRAM if needed [Gru20]. Most importantly, requests to data that is present in the caches is served faster than if it were not.

A visualization of the aforementioned components can be found in fig. 1.

2.2 OUT-OF-ORDER EXECUTION

As the name implies, out-of-order execution refers to the idea of executing instructions in a different order than the one in which they are given [Gru20]. With multiple execution units that run in parallel (as described in sec. 2.1), CPUs can take advantage of mutually independent instructions and execute them at the same time.

The basic realization of this concept is provided by Tomasulo’s algorithm [Tom67]. It introduces two components, the first of which is the reservation station, which collects the operands of instructions until they are ready to be executed by the execution units. Crucially, the corresponding execution units do not need to wait until all operands are present and can instead compute other instructions. The second component of Tomasulo’s algorithm is the Common Data Bus (CDB), which connects all reservation stations and execution units. Whenever a result is computed by an execution unit, the result is broadcast onto the CDB and thus made available to all reservation stations that are



FIGURE 1: Simplified overview an Intel Skylake CPU [Gru20, fig. 2.1]. For the memory subsystem, detailed knowledge of the load and store buffers, as well as the TLBs, is not required. The same applies to the allocation queue of the frontend.

waiting for it. This important step ensures that results are not written to registers first, just to be read again by other instructions that need them as operands.

Initially, according to Tomasulo, each set of execution units needed its own reservation station, however, more modern implementations by Intel use a single unified reservation station, called scheduler, that handles all types of instructions, rather than just one [Tom67] [Wik] [Gru20]. This can also be seen in fig. 1.

2.3 SPECULATIVE EXECUTION

Speculative execution allows a CPU to predict the outcome of comparisons and other branch instructions. This prevents stalls when waiting for the instruction that determines which branch is taken to finish. Similarly to out-of-order execution, rollbacks are needed in some cases. However, in addition to exceptions, they also occur if a branch was mispredicted. [Gru20]

To predict the outcome of branch instructions, CPUs include a branch prediction unit (BPU) [Gru20]. While available in different configurations, many modern CPUs record the most recent outcomes of a branch with a counter [Gru20] that is either incremented if the branch is taken, or decremented.

2.4 MELTDOWN AND SPECTRE

Meltdown [Lip+18] and Spectre [Koc+19] abuse out-of-order and speculative execution to leak data from memory addresses that are normally inaccessible to the attacker over the caches of the CPU.

2.4.1 MELTDOWN

On a high level, Meltdown [Lip+18] works by forcing exceptions when reading data inaccessible to the attacker and transiently encoding this data into the cache to retrieve it once the rollback has completed. What enables Meltdown is a small time window between an invalid memory access and the raising of an exception [Lip+18].

The Meltdown-US-L1 variant of Meltdown [Lip+18] to which the presented emulator is vulnerable works by accessing a memory address for which the attacker has no permission. Firstly, the attacker allocates an oracle array and ensures none of its entries are present in the cache. Upon loading the contents of the inaccessible memory location into a register, the attacker uses the value to access the array at a specific offset. When the rollback caused by the access violation has completed, the attacker measures the access time to each of the array entries to determine which one has been accessed. It is important that the data to be stolen is currently cached. While there are numerous other variants of Meltdown that differ from the basic variant mainly by how they force an exception and from which microarchitectural buffer they leak data, these types of attacks are out of scope.

2.4.2 SPECTRE

Spectre [Koc+19], on the other hand, relies on the CPU mispredicting a branch and transiently executing instructions that are not part of the correct execution path. This misprediction in a victim process can be induced by maliciously configuring the BPU (sec. 2.3). Depending on the instructions that are wrongfully executed by the victim, traces may later be found in the processor’s cache. Similarly to Meltdown, different variants of Spectre exist. The presented emulator is vulnerable to variant 1 of Spectre, which takes advantage of the CPU mispredicting the outcome of comparison instructions.

2.4.3 MITIGATIONS

The mitigations for Meltdown are available in both software and hardware, some of which are implemented in the emulator to allow users to experiment with them and determine their effectiveness. A first, rather simple, mitigation for Meltdown is to disable out-of-order execution [Lip+18], which would completely prevent an attacker from encoding the normally inaccessible data into the cache. Later revisions of Intel’s architectures introduced further mitigations. Although undocumented by Intel, researchers suspect the processor still performs the illegal read, but zeros out the data that is given to dependent instructions before raising an exception [Can+20]. The emulator implements this mitigation, which may be enabled by the user. Other mitigations deployed by operating systems, such as KTPI (or KAISER) [Lip+18], are highly effective, but out of scope, as there exists no operating system.

Unlike Meltdown, Spectre appears to be a design flaw. While one might argue that transiently computing on real values after it is already known an exception will occur is a bug, the behavior abused by Spectre is a direct consequence of speculative execution (sec. 2.3). As a result, an effective yet questionable mitigation is to simply disable speculative execution [Koc+19]. Alternatively, Intel recommends potential victim’s of Spectre v1 to use an “lfence” instruction where appropriate, which ensures prior load instructions retire before continuing, thus effectively disabling speculative execution for certain parts of an application [Int18]. A fence instruction is provided as part of the emulator’s instruction set.

An additional mitigation that works against both Meltdown and Spectre is to flush the entire cache after a rollback. While highly inefficient, it does prevent the retrieval of otherwise inaccessible data after the transient execution phase of the attacks. This mitigation can be realized using the emulator by a sequence of instructions (microprogram) that are executed after each rollback.

3 SPECIFICATION OF OUR TASK

todo

4 CPU EMULATOR/ BACKEND

Melina Hoffmann

In this chapter, we introduce the backend of our emulator program. It contains the elements of our program that emulate actual CPU components. Our emulator is based on information about modern real life CPUs, especially the x86 Intel Skylake architecture [Wik].

In our program, we model the distinct components of a real life CPU with a modular setup making use of the object oriented functionalities of Python3. Breaking up the source code into individual CPU components also makes it easier to understand and maintain. Additionally, we have made simplifications and modifications in comparison to a real life CPU, so the emulator is as clear and easy to understand as possible while still implementing an actual out of order execution and providing the necessary functionality for Meltdown and Spectre attacks.

In this chapter we firstly introduce the components of our CPU emulator, how they work (together) and which part of a real life CPU they emulate sec. 4.1. Then we explain how our emulator provides out of order execution and how it may differ from the general Tomasulo algorithm introduced in sec. 2.2. Subsequently, we show how we implemented rollbacks and exception handling, especially with regards to how our implementation allows for Meltdown and Spectre attacks sec. 4.3. Then we give an overview over our instruction set architecture sec. 4.4. Lastly we show how our emulator can be adapted for different demonstrations and attacks without changing the source code via a config file sec. 4.5.

4.1 CPU COMPONENTS AND OUR EQUIVALENTS/ MODELS (10-11 PAGES)

todo

4.1.1 CPU

The primary purpose of the CPU module is to allow all other components to work together. That is, the CPU initializes all other components and provides interfaces to their functions, which allows the GUI to visualize the current state. In addition, the CPU provides functions which load user programs and a tick function that is called each cycle and does the following: instructions from the instruction queue are fetched from the frontend and forwarded to the execution engine, until either no more slots in the unified reservation station are available, or the instruction queue is empty. It then calls the tick function of the execution engine. In case of a rollback, the instruction queue maintained by the frontend is flushed. If the rollback was caused by a faulting load instruction, execution is resumed at the next instruction (as described in sec. 4.3). In case a branch was mispredicted, the frontend is notified and refills the instruction queue accordingly. If configured, corresponding

microprograms are sent to the instruction queue. Lastly, a snapshot of the current state of the CPU is taken.

The second purpose of the CPU is to provide the snapshot functionality which allows a user of the emulator to step back to previous cycles. The snapshot list is simply a list which grows at each cycle, where each entry is a deep copy of the CPU instance. However, this would imply that each snapshot maintains a full list of snapshots, where each entry contains a list of snapshots, and so on. To combat this, there is only a single global list of snapshots, and each snapshot entry maintains a reference and an index to its own entry. As a result, the default deep copy function is adjusted accordingly. Due to the low complexity of the programs we expect users to run, at the moment, no maximum number of available snapshots is configured.

4.1.2 INSTRUCTIONS AND PARSER

todo

4.1.3 DATA REPRESENTATION

todo

4.1.4 CPU FRONTEND

In modern CPUs, the CPU frontend provides an interface between code and execution engine. It contains components to fetch and decode the instructions from a cache and supply them to the CPU in a queue. It is also involved in speculative execution by predicting the result of conditional jumps and supplying instructions to the execution engine accordingly. [Wik]

Melina Hoffmann

In our emulator we simplify the components and procedures. We also separate the branch prediction unit (BPU), that manages and predicts the results of conditional jumps, from the rest of the frontend. This makes our code easier to understand and makes potential future changes or additions to the BPU more convenient.

BRANCH PREDICTION UNIT (BPU)

The BPU of modern CPUs plays a vital role in enabling speculative execution. It stores information about previously executed conditional branch instructions and predicts the outcome of future branch instructions accordingly. The rest of the CPU can then speculatively execute further instructions at an address based on the predicted outcome of the branch instead of stalling until the branch instruction is processed by the execution engine. If the prediction was false, a rollback is performed on the speculatively executed instruction sec. 4.3. If the prediction was true, the execution is overall faster than without speculative execution.

Detailed information about the BPU of our modern base CPU is not widely available [Wik]. In general, the Gshare BPU of modern CPU consists of multiple components. It contains a branch target buffer (BTB) that holds the predicted target addresses for conditional branches. It also uses a

pattern history table (PTH), that can be implemented as a 2-bit-saturating counter, and a global history register (GHT) to predict whether a conditional jump will be taken or not. [SCA]

In our emulator, we forgo the BTB entirely. Real life CPUs benefit from stored target addresses since they have to expensively decode each branch instruction before they can work with the target address [Wik]. In our emulator, the parser decodes the jump labels from the assembler code and directly provides them to the frontend, so storing them in an additional buffer is unnecessary sec. 4.1.2. We further forgo the GHT because it is not strictly necessary to execute a Meltdown or Spectre attack. Additionally, our emulator and its behaviour are easier to understand and predict without it, which is important when implementing microarchitectural attacks for didactic purposes. The BPU of our emulator only consists of a PHT, which is enough for simple Spectre-PHT variants [reference-eval-spectre].

The default PHT used in our emulator holds an array called counter of configurable length to store several predictions. The instructions are assigned to different prediction slots by the last bits of their index in the instruction list. For each of the slots, the prediction can take the four different values from zero to three, where zero and one indicate that the branch will probably not be taken and two and three indicate that the branch is likely to be taken.

The source code for our emulator also contains a more simple BPU with only one slot for all instructions. Since the number of slots in the default BPU is freely configurable by the user, this simple BPU is now obsolete.

When the BPU is updated with an actual branch outcome from the execution engine, the prediction in the PHT is updated by a 2-bit counter. This means, that if the prediction was right, the counter remains at or updates to zero (strongly not taken) or three (strongly taken) respectively. If the prediction counter is at zero but the branch is actually taken, the counter is updated to one (weakly not taken). If it is at a one when the branch is taken, it is directly updated to three. The counter behaves similarly when it has the value two or three and the branch is not actually taken.

INSTRUCTION QUEUE

In a real life CPU, the overall purpose of the frontend is to provide the execution unit with a steady stream of instructions so the backend is busy as much as possible and therefore efficient. In a modern x86 CPU, the frontend has to fetch x86 macro-instructions from a cache and decode, optimize and queue them repeatedly to provide the backend with a queue of μ -instructions ready for issuing in the execution engine. [Wik]

In our emulator, except for the BPU, the functionality of the CPU frontend is bundled in frontend.py. It is significantly simplified compared to a real life x86 CPU, especially since we use only one type of instructions instead of distinguishing between macro- and μ -instructions sec. 4.4. They are already provided as a list by the parser, which renders the decoding and optimization steps in the frontend unnecessary sec. 4.1.2.

The main task of our frontend is to act as interface between instruction list provided by the parser and the execution engine sec. 4.1.6. It provides and manages the instruction queue, which holds the instructions that the execution engine should issue next. Conditional branches with their respective

BPU predictions are taken into account when filling the queue. This enables speculative execution which is needed for Spectre attacks sec. 2.4.

The central component of our emulated frontend is the instruction queue. In our version, it does not only hold the instructions themselves, but also for every instruction in the queue, the respective index in the instruction list is stored. For branch instructions it also holds the respective branch prediction from the BPU at the time that the instruction was added to the queue. This additional information is needed by the execution engine to handle mispredictions and other exceptions sec. 4.1.6.

When adding instructions to the queue, the frontend selects them from the instruction list, adds the additional information for the execution engine and places them into the instruction queue until the queue's maximum capacity is reached. The frontend maintains a program counter (pc) that points to the next instruction in the list that should be added to the queue. When the frontend encounters a branch instruction and the branch is predicted to be taken, the frontend adjusts the pc to resume adding instructions at the branch target. If a branch was mispredicted, the frontend provides a special function to reset the pc and refill the instruction queue with the correct instructions.

Additionally, the frontend provides a function to add a μ -program to the queue. It consists of a list of instructions separate from the parser instruction list. When adding the μ -program to the queue, the frontend may exceed the maximum queue capacity. This functionality can be used to implement mitigations against microarchitectural attacks, e.g. by adding a μ -program as part of the exception handling after an illegal load [**reference-eval-mitigations**].

The frontend provides interfaces to both read and take instructions from the queue. It also provides a function that combines taking an instruction from the queue and refilling it, . Additionally, the frontend has an interface for flushing the whole queue at once without taking the instructions from the queue. This can be used when demonstrating mitigations against microarchitectural attacks [**reference-eval-mitigations**].

The frontend provides further basic interfaces, e.g. for reading the size of the instruction queue and reading and setting the pc. These are used by the other components during regular execution, e.g. when issuing instructions to the execution engine, but also to reset the queue to a certain point in the program after an exception has occurred sec. 4.3. Since our emulator only executes one program at a time, the other components can check via another interface whether the frontend has reached the end of the program.

4.1.5 MEMORY

todo

4.1.6 EXECUTION ENGINE

todo

4.2 OUT OF ORDER EXECUTION (2 PAGES)

Our emulator implements out-of-order execution. This allows transient execution of instructions before the fault handling of previous instructions is finalized, which is essential for Meltdown type attacks sec. 2.4. Our version of out-of-order execution is based on Tomasulos algorithm sec. 2.2. Since the goal of our out-of-order execution is to enable Meltdown attacks and make them easy to understand, not quick performance, we use a basic/ simple version of Tomasulos algorithm. In our emulator, the components necessary for Tomasulos algorithm are located in the Execution Engine sec. 2.2, sec. 4.1.6. Below, we provide a detailed look at our version of out-of-order execution and the components involved in its implementation.

4.2.1 ISSUING INSTRUCTIONS

Since we implement out-of-order execution according to Tomasulos algorithm, our Execution Engine does not try to execute instructions directly when it receives them from the frontend sec. 4.1.4, sec. 4.1.6. Instead, it issues them to the Reservation Station where multiple instructions can wait until all their operands are ready. If all operands of an instruction are ready, the Execution Engine can execute it. This does not generally happen in the order of instructions provided by the program, but will always lead to the right effect in that each instruction is executed with the right operands as determined by the program.

The instructions are provided by the frontend in program order sec. 4.1.4. The Execution Engine issues them into the Reservation Station, if it is not yet fully occupied. The Reservation Station is modelled as a list of slots which can each hold an instruction together with additional information about the instruction. It is unified in that each spot in the list can hold slots for all types of instructions.

To issue an instruction, the Execution Engine creates a slot object that fits the type of the instruction and puts it into the Reservation Station. Besides the instruction itself, it holds additional information, including a list of the operands. While immediate operands can be directly converted to a Word, register operands have to be resolved during the issuing process.

The registers are modelled by a list in which each entry can either be a Word value or the ID of the slot in the Reservation Station which produces the next register value. Since the instructions are issued in program order, this reflects the expected register state at the point of issuing the current instruction if the program was executed in order. The only difference being, that the results of yet unexecuted instructions are being represented by the respective slot ID. sec. 4.1.6 To resolve the register operands, the current content of the respective register is added to the operand list, so the operand list can contain both Words and slotIDs. As described below, slotIDs in the operand list will be replaced by the result of the instruction producing the value when it finishes executing.

Resolving the register operands this way ensures that data dependencies between instructions that use the same registers are adhered to. To increase performance, real life CPUs practice register renaming by assigning ISA level registers to different μ -architectural registers in order to further eliminate data dependency hazards [Gru20]. Since we do not differentiate between the ISA and μ -architectural level sec. 4.4 and aim to keep our emulator easy to comprehend, we do not implement register renaming.

Once the slot with the new instruction is placed into the Reservation Station, if the instruction has a target register for its result, the slotID of the instruction is put into this target register. This ensures, that when the next instruction is issued, the register state again represents the expected register state if the instructions were executed in-order. Note that placing the slotID into the target register cannot only overwrite a Word value but also a slotID, if the previous instruction that uses this register as its target register is not yet fully executed. This is not a problem, since every instruction that may need the result of the respective instruction for the previous slotID as an operand, already holds this slotID in their own operands list. It will be notified of the result when it is ready, regardless of whether the slotID is still present in the register or not.

4.2.2 EXECUTING INSTRUCTIONS

In a real life CPU/ basic Tomasulo, when all operands of an instruction in the Reservation Station are ready, it is transferred to a free execution unit and executed. In our emulator, execution of the instructions is triggered by the tick function of the Execution Engine, which is executed once per CPU cycle. We do not model a finite number of execution units as separate components sec. 4.1.6. Instead, the tick function goes through the occupied slots in the Reservation Station and tries to execute each instruction by calling the tick_execute function of its respective slot. This follows the order of the slots in the Reservation Station, regardless of when the instruction in each slot was issued, i.e. regardless of their order in the program.

If the operands of the current instruction are not ready yet, i.e. there are still slotIDs in the operand list, the instruction is skipped. To mimic the latency of real world execution units and memory accesses, each instruction type waits a specific amount of ticks after all operands are ready until producing its result. These waiting instructions are also skipped.

Once the instruction is executed and produces a result, i.e. all operands are available and the wait time is over, according to Tomasulos algorithm this result has to be broadcasted via the CDB to the other slots and the registers. In our emulator, the CDB is modelled by the _notify_result function. It goes through all registers and replaces all occurrences of the slotID of the instruction with the result it just produced. It also notifies all occupied slots of the result together with the slotID of the instruction which produced it, so they can replace the slotID in their operands list if it occurs. If a result is produced like this, the tick function returns without executing further slots. This mimicks that a real life CDB can only broadcast one result each cycle. It has the side effect that instructions do not necessarily execute in the same number of ticks, depending on where they sit in the list of slots. The tick function also returns before all slots have been executed if the instruction in a slot retires, in order to properly handle potentially faulting instructions sec. 4.3.

4.2.3 MEMORY HAZARDS

As described above we handle data dependencies between instructions that use the same registers by using slotIDs as placeholders for as yet uncomputed results. We also need to handle data dependencies between memory accesses. For this, each slot that contains a memory instruction also holds a set of slotIDs of other memory instructions that potentially lead to a memory hazard together

with the current instruction. The memory instruction is only executed when all other instructions from its list of potential hazards have finished executing.

This list is filled when the memory address the instruction will access is computed, beforehand it is set to the placeholder value *none*. To fill the list, the `*_tick_execute*` function of the slot goes through its list that contains all in-flight memory instructions that precede the current instruction in program order, and includes them if they access the same address. If there is a previously issued memory instructions for which the memory address is not yet available, the instruction waits until the hazard list can be completed.

By adding all instructions to the hazard list that access the same memory address, we potentially generate false positives in the case that two *load* operations read from the same memory address right one after the other. Since efficiency is not our priority, we accept this in order to keep our emulator simple. Additionally, to simplify fault handling, *store* instructions wait until all other possibly faulting instructions are executed.

4.2.4 FENCE INSTRUCTION

The *fence* instruction is a special instruction in that it does not produce a result or a lasting side effect in the other components of the emulator. It creates a fixed point in the execution of the program sec. 4.4. It holds a list of all instructions that were already in the Reservation Station when itself was issued. Similar to the memory instructions, it waits for all the instructions in the list to be retired before executing itself. Additionally, no other instructions can be issued to the Reservation Station while it contains a fence instruction. This effectively suspends the out-of-order execution with regards to the fence instruction, since all previous instructions in the program order are executed before the fence and all following instructions are only issued, and therefore executed, after the fence was executed. This can be used to model mitigations against μ -architectural attacks [ref_mitigations_eval].

4.3 EXCEPTIONS AND ROLLBACKS (2-3 PAGES)

todo

4.4 ISA (2 PAGES)

Real life Intel x86 CPUs differentiate between two types of instructions or operations. Macro-operations refer to the relatively easily human readable and convenient but complex instructions that are described by the x86 ISA. Their length differs between the instructions. Internally, in the execution units, the CPU works on μ -operations, which are small operations of a fixed length. One macro-operation contains one or multiple μ -operations. The CPU frontend has to decode the macro-operations into μ -operations in an expensive multi step process. sources: [Wik], <https://en.wikipedia.org/wiki/macro-operation>, <https://en.wikipedia.org/wiki/micro-operation>

Melina Hoffmann

Our CPU emulator only uses one type of instructions. They are directly read from our assembler code by the parser and passed to the execution engine without further decoding, splitting or replacing sec. 4.1.2, sec. 4.1.4. To show basic Meltdown and Spectre variants, we do not need overly complex instructions, e.g. instructions that contain multiple memory accesses in one or that are used to perform encryption in hardware [ref_evaluation_meltdown], [ref_evaluation_spectre]. Basic arithmetic operations, memory accesses, branches and a few special operations are sufficient for the demonstrated attacks and are both easy to implement as single instructions and to use in assembler code that should be well understood by the author. Using the same operations throughout the emulator also makes the visualization more clear and easier to follow, e.g. when the same operations appear, one after the other, in the visualization of the assembler code, the instruction queue and the reservation stations [ref_ui].

4.4.1 DEFAULT INSTRUCTION SET

In order that our CPU emulator can recognize and work with an instruction, it has to be registered with the parser sec. 4.1.2. In our default setting default, we register a basic set of instructions with the parser so students can start writing assembler code and using the emulator right away. This basic instruction set is also used in our example programs in [ref_UI].

Our relatively small instruction set is based on a subset of the RISC-V ISA. It offers a selection of instructions that is sufficient to implement Meltdown and Spectre attacks as well as other small assembler programs while still being of a manageable size so students can start to write assembler code quickly without spending much time to get to know our ISA. The syntax of the assembler representation is also based on RISC-V (as introduced in the “RISC-V Assembly Programmer’s Handbook” chapter of the RISC-V ISA) [ref_RISC-V]. If needed, students can add further instructions by registering them with the parser sec. 4.1.2.

In the following subchapters we introduce the instructions of our default ISA. They are grouped according to their respective instruction type in the emulator except for the special instructions which are grouped together sec. 4.1.2. All default instructions are summarized in the appendix into a quick reference sheet [ref_appendix].

ARITHMETIC AND LOGICAL INSTRUCTIONS WITHOUT IMMEDIATE

These are basic arithmetic and logical instructions that operate solely on register values, i.e. both source operands and the destination operand reference registers. For simplicity, we write, for example, Reg1 when referring to the value read from or stored in the register referenced by the first register operand.

Each of these default operations uses the respective python standard operator on our Word class to compute the result, except for the right shifts. For the logical and the arithmetic right shift, the python standard right shift operator is used on the unsigned and the signed version of the register value respectively. When returning the result as a Word, it is truncated to the maximal word length by a modulo operation, if necessary sec. 4.1.3. This means, that any potential carry bits or overflows are effectively ignored.

Arithmetic and Logical Instructions without Immediate		
Instr. Name	Operators	Description
add	Reg1, Reg2, Reg3	Reg1 := Reg2 + Reg3
sub	Reg1, Reg2, Reg3	Reg1 := Reg2 – Reg3
sll	Reg1, Reg2, Reg3	Reg1 := Reg2 << Reg3
srl	Reg1, Reg2, Reg3	Reg1 := Reg2 >> Reg3 logical
sra	Reg1, Reg2, Reg3	Reg1 := Reg2 >> Reg3 arithmetical
xor	Reg1, Reg2, Reg3	Reg1 := Reg2 xor Reg3
or	Reg1, Reg2, Reg3	Reg1 := Reg2 or Reg3
and	Reg1, Reg2, Reg3	Reg1 := Reg2 and Reg3

ARITHMETIC AND LOGICAL INSTRUCTIONS WITH IMMEDIATE

These are basically the same instructions as in sec. 4.4.1. The main difference is, that the second source register is replaced by an immediate operand which is set directly in the Assembler code. This immediate is used as the value of a Word in the execution engine, so it is truncated by a modulo operation to be in the appropriate range sec. 4.1.3, sec. 4.1.6.

Arithmetic and Logical Instructions with Immediate		
Instr. Name	Operators	Description
addi	Reg1, Reg2, Imm	Reg1 := Reg2 + Imm
subi	Reg1, Reg2, Imm	Reg1 := Reg2 – Imm
slli	Reg1, Reg2, Imm	Reg1 := Reg2 << Imm
srl	Reg1, Reg2, Imm	Reg1 := Reg2 >> Imm logical
srai	Reg1, Reg2, Imm	Reg1 := Reg2 >> Imm arithmetical
xori	Reg1, Reg2, Imm	Reg1 := Reg2 xor Imm
ori	Reg1, Reg2, Imm	Reg1 := Reg2 or Imm
andi	Reg1, Reg2, Imm	Reg1 := Reg2 and Imm

MEMORY INSTRUCTIONS

These instructions provide basic interactions with the emulated memory sec. 4.1.5. Load and store instructions exist in two versions, one that operates on Word length data chunks, for convenience, and one that operates on Byte length data chunks, for the fine granular access needed in micro architectural attacks. The flush instruction flushes the cache line for the given address sec. 4.1.5. The address is calculated in the same way for all memory instructions: $\text{addr} := \text{Reg2} + \text{Imm}$, and $\text{addr} := \text{Reg} + \text{Imm}$ for the flush instruction respectively.

Memory Instructions		
Instr. Name	Operators	Description
lw	Reg1, Reg2, Imm	Reg1 := Mem_word[addr]
lb	Reg1, Reg2, Imm	Reg1 := Mem_byte[addr]
sw	Reg1, Reg2, Imm	Mem_word[addr] := Reg1
sb	Reg1, Reg2, Imm	Mem_byte[addr] := Reg1
flush	Reg, Imm	flush cache line of addr

BRANCH INSTRUCTIONS

All branch instructions compare the values of two source registers. If the comparison evaluates to true, the execution of the program is resumed at the given label in the assembler code. If it evaluates to false, the next instruction in the program is executed. Depending on the instruction, the register values are interpreted as signed or unsigned integers. Labels in the assembler code are automatically resolved by the parser sec. 4.1.2, [rev_eval_and_example_code].

Branch Instructions		
Instr. Name	Operators	Description
beq	Reg1, Reg2, Label	jump to Label if Reg1==Reg2
bne	Reg1, Reg2, Label	jump to Label if Reg1!=Reg2
bltu	Reg1, Reg2, Label	jump to Label if u(Reg1)<u(Reg2)
bleu	Reg1, Reg2, Label	jump to Label if u(Reg1)<=u(Reg2)
bgtu	Reg1, Reg2, Label	jump to Label if u(Reg1)>u(Reg2)
bgeu	Reg1, Reg2, Label	jump to Label if u(Reg1)>=u(Reg2)
blts	Reg1, Reg2, Label	jump to Label if s(Reg1)<s(Reg2)
bles	Reg1, Reg2, Label	jump to Label if s(Reg1)<=s(Reg2)
bgt	Reg1, Reg2, Label	jump to Label if s(Reg1)>s(Reg2)
bges	Reg1, Reg2, Label	jump to Label if s(Reg1)>=s(Reg2)

SPECIAL INSTRUCTIONS

Rdtsc acts like a basic timing instruction. It returns the number of ticks the execution unit has executed so far in the given register.

The fence instruction acts as a fixed point in the out of order execution. All instructions that are already issued in the execution unit at the point of issuing the fence instruction are executed before the fence is executed. No new instructions are issued before the execution of the fence instruction is complete.

Special Instructions		
Instr. Name	Operators	Description
rdtsc	Reg	Reg:=cyclecount
fence	-	add execution fixpoint at this code position

4.5 CONFIG FILES (1 PAGE)

todo

5 USER INTERFACE AND USAGE

todo

6 EVALUATION

todo

7 CONCLUSION

todo

8 CATCHY LAB TITLE

8.1 ABSTRACT

8.2 INTRODUCTION (1 PAGE)

```
1 general task/ goal
2     CPU emulator with actual out of order execution, maybe speculative
3     execution and comprehensible implementation and documentation
4     implement different p-architectural features like out-of-order exe.,
5     speculative exe. b.c. we want to observe and demonstrate p-arch.
6     attacks
7 brief context of the task/ goal (e.g. a sentence on Meltdown and why it is
    important to understand it)
8     why is it important to have the emulator: so far no possibility to
9     observe microarchitectural attacks in real life
10    rauscharme Umgebung, die macht was man erwartet
11 structure of the report
```

8.3 BRIEF THEORETICAL BACKGROUND (2-3 PAGES)

```
1 premise according to Felix: reader knows SCA lecture -> brief recaps/
2     reminders to reference back to from the other chapters
3
4 ### really brief introduction to CPU
5
6     multiple components
7         maybe mostly via picture
8
9     how they work together
10
11     maybe data-flow from instruction to result
12
13     a lot of hypothticals due to trade secrets
14
15     suggested literature:
16         maybe Gruss Diss. and other works on cbsca
17         maybe textbooks
18         maybe CPU wiki for pictures
19
20 ### out-of-order execution
```

```

20
21     a bit more in depth b.c. this is not content of the sca lecture
22
23     maybe why do we need it/ advantages
24
25     brief explanation of Tomasulo and how it works
26
27     suggested literature:
28         original Tomasulo paper from 1965
29         maybe a more recent/ didactically edited explanation (is this in
30         textbooks?)
31
32     also something about speculative execution
33
34 ### Meltdown and Spectre
35
36     brief introduction to Meltdown (and Spectre) in general
37
38     slightly more in-depth introduction to the Meltdown attack we picked
39     and want to run in our emulator
40
41     highlight the relevant parts of the CPU, maybe mention how they
42     interact in Meltdown (and Spectre)
43
44     mitigations and maybe how successfull they are
45
46     suggested literature:
47         ggf. Dissertation Gruss if we mention basic cbsca
48         Gruss et al. papers on Meltdown and Spectre e.g. https://gruss.cc/
49         files/meltdown.pdf or https://gruss.cc/files/meltdown\_cacm.pdf
50         Canella, Gruss et al. on mititgations https://gruss.cc/files/
51         transient-attacks.pdf

```

8.4 MAYBE SPECIFICATION OF THE TASK (1 PAGE)

```

1   rauswerfen, in Introduction mit übernehmen
2
3   emulator to execute and teach Meltdown
4       who is the target audience`
5       which Meltdown attacks specifically
6       etc. further concretisations
7
8   are there existing solutions/ related works?
9       Felix' old emulator? citable?
10      suggested literature:
11          todo

```

8.5 CPU EMULATOR/ BACKEND (17-18 PAGES)

```

1 maybe general info e.g. that we wrote it in python, if we used special
   tools/ libraries
2 for everything we implemented:
3     which part of a real life cpu does this model/ how is this done in
   real life cpus
4     briefly how does it work
5     why did we choose to model it like we do
6         nice code structure/ code easy to understand by students
7         some features more or less relevant for meltdown
8         etc.
9     maybe what did we leave out
10    challenges?
11
12 ### CPU Components and our equivalents/ models (10-11 pages)
13
14     modular setup based on real life CPUs and nice coding conventions
15
16     #### CPU
17
18         contains/ controls the rest
19
20         preparation (loading the program, init the rest)
21
22         ticks
23
24         coordinate rollbacks?
25
26     #### Instructions and Parser
27
28         instructions
29
30         parser
31
32     #### Data representation
33
34         how we model and handle data
35
36         byte, word
37
38     #### CPU frontend
39
40         Branch Prediction Unit (BPU)
41
42         Instruction Queue
43
44     #### memory
45

```

```

46         no virtual addresses
47
48         mmu
49
50         how do we store data
51
52         how do we model data access (i.p. wrt Meltdown)
53
54         cache
55
56         what kinds of caches can we represent
57
58         how do we model if something is cached and the access times (i.
59         p. wrt Meltdown)
60
61     ##### Execution Engine
62
63         Reservation Station/ Slots
64
65         unlimited execution units
66
67     ### Out of Order Execution (2 pages)
68
69         our version of out-of-order execution/ Tomasulo
70         how much details on Tomasulo? our choice
71
72         where in our program do we implement which components
73
74         what did we leave out/ do differently
75
76     ### Exceptions and Rollbacks (2-3 pages)
77
78         in particular wrt making Meltdown possible
79
80         general goal/ concept of rolling back after a misprediction or an
81         exception
82
83         snapshots and interaction of the CPU components
84         components that are fully reset after a rollback: cache, BPU
85
86     ### ISA (2 pages)
87
88         overview of ISA
89
90         p-instr. vs. ISA
91         mixture of both in one: only p-code, but more abstract than in real
92         life
93
94         do not need to think about page faults etc. anyways, do not need
95         overly complex instructions

```

```

91
92     reasoning behind the choice of instructions (manageable size and
          instructions (e.g. no divide by zero) balanced with functionality
          particularly wrt Meltdown)
93
94
95     ### config files (1 page)
96
97     what can be configured without changing the source code
98         why these variables? relevant for Meltdown?
99         why did we choose which default values?

```

8.6 OUR VISUALISATION AND USAGE/ FRONTEND (10 PAGES)

```

1   ggf. gemäß Anmerkung von Lenni umsortieren: die idee das UI nicht mit in
          die 'unsere designentscheidungen' zu nehmen, sondern im prinzip so als
          Manual abzukapseln finde ich eigentlich ganz gut. Müssen wir aber dann
          mal in der Praxis schauen. Soll ja keine didaktischen begründungen das
          manual zu sehr aufblähen, vllt wird das sinnvoll, dann einen teil des
          UI im "backend" kapitel einzubauen, und dann wirklich ein cleanes
          manual kapitel zu haben
2
3   zwei Strukturierungsmöglichkeiten:
4       z.B. alphabetisch
5       z.B. nach einzelnen Angriffen etc. aufgeteilt und aufeinander aufbauend
6   erklären was technisch passiert nachdem man gezeigt hat wie es aussieht
7   man kann auch explizit sagen, "ich gehe davon aus, dass du GDB kannst,
          daran ist das angelehnt, hier sind die Unterschiede"
8
9   ### general concept
10
11     goal: UI for the emulator with visualisation of CPU/ memory components
          and their contents
12
13     in terminal
14         maybe comparison to existing analysis tool/ debugger gdb
15
16     triggers/ controls the actual emulator
17         overview features, e.g. breakpoints, step-by-step and stepback
18
19     maybe which tools/ libraries were used?
20
21     ### features in more detail and their didactic purpose
22
23     breakpoints, step-by-step, stepback and more
24
25     challenges/ design choices during the implementation

```

8.7 DEMONSTRATION/ EVALUATION (7 PAGES)

```

1  what kind of system do we need/ did we use to run this?
2      which python Version?
3      which program version ? git commit
4      other dependencies?
5
6  ### general demonstration
7
8      brief example program showing all the features in a "normal" execution,
        e.g. adding stuff
9
10 ### Meltdown und Spectre demonstration
11
12     #### Example Program Meltdown
13
14         show example program
15
16         maybe compare to example program for real life architecture from
        SCA or literature (Gruss) if available
17
18         explain which Meltdown variant it implements
19
20         briefly highlight which components (we expect to) interact to make
        it work
21
22         how well does it work?
23
24     #### maybe example program spectre
25
26         same as Meltdown
27
28     #### mitigations
29
30         it is known which mitigations exist, here is what we have in our
        emulator:
31
32         what is possible in our program as is
33             planned:    cache flush: microcode -> config file
34                       mfence in assembler (normally in compiler)
35                       aslr directly in program -> config (es gibt ja auch
        mitigations, die keine echte mitigation sind; nice to have -> könnte
        demonstrieren dass es nicht der Fall ist; war eh schon einige Jahre vor
        Meltdown vorhanden/ in Gebrauch; KSLR brachen kann man auch als
        Angriff verkaufen)
36
        flush IQ -> passiert eh schon, ist das überhaupt
        eine echte mitigation?
37
        disable speculation (nice to have, lassen wir weg)
        -> config

```

```

38         out of order -> in config RS mit nur einem Slot
39
40         is our meltdown/ spectre variant still possible?
41         ggf. how does this affect the performance?
42         vorsichtig sein, dass man dann auch die richtige Frage für die
Antwort stellt
43         in real life (already in background)
44         in our program
45
46         what would be the necessary steps/ changes to the program for
further mitigations
47         compare to changes in hardware by the manufacturers

```

8.8 CONCLUSION (1 PAGE)

```

1  recap goals
2  main goal of this chapter: to which extend did we reach our goals?
3
4  did we reach the goal of implementing a CPU emulator where a user can
perform a Meltdown attack
5
6  how many Meltdown attacks are possible?
7
8  is Spectre possible?
9
10 mitigations
11     which mitigations did we implement
12     is this a good amount/ sample of real world mitigations or do we miss
important ones?
13     how well do they perform (also in comparison to how they perform in the
real world)
14
15 value to students:
16     how easy to use and convenient do we think our program is?
17     do we think this will be a good tool for teaching?
18     how accessible is it wrt different host architectures?
19     1 überzeugter Satz: wir haben sehr gute Arbeit geleistet und das
Werkzeug ist sehr gut. wenig hin schreiben, damit keiner fragt, ob wir
das nicht weiter evaluieren müssten. nicht weit aus dem Fenster mit
Behauptungen lehnen sondern klar und kurz unsere Meinung formulieren
20
21 further work
22     more (detailed/ realistic) functionality for more Meltdown and Spectre
variants
23     more elaborate BPU with btb (and ghr) for more spectre variants
24     more mitigations
25     maybe nice to have wrt to the visualisation/ general UI functionality/
ISA?

```

REFERENCES

- [Can+20] CANELLA, Claudio et al.: “KASLR: Break it, fix it, repeat”. In: *Proceedings of the 15th ACM Asia Conference on Computer and Communications Security*. 2020, pp. 481–493.
- [Gru20] GRUSS, Daniel: “Transient-Execution Attacks”. 2020. URL: <https://gruss.cc/files/habil.pdf> (visited on 03/11/2022).
- [Int18] INTEL: *Intel Analysis of Speculative Execution Side Channels*. 2018. URL: <https://newsroom.intel.com/wp-content/uploads/sites/11/2018/01/Intel-Analysis-of-Speculative-Execution-Side-Channels.pdf> (visited on 03/11/2022).
- [Koc+19] KOCHER, Paul et al.: “Spectre attacks: Exploiting speculative execution”. In: *2019 IEEE Symposium on Security and Privacy (SP)*. 2019.
- [Lip+18] LIPP, Moritz et al.: “Meltdown: Reading Kernel Memory from User Space”. In: *27th USENIX Security Symposium (USENIX Security 18)*. 2018.
- [Tom67] TOMASULO, Robert M: “An efficient algorithm for exploiting multiple arithmetic units”. In: *IBM Journal of research and Development* 11.1 (1967), pp. 25–33.
- [Wik] WIKICHIP: *Skylake Microarchitecture*. URL: https://en.wikichip.org/wiki/intel/microarchitectures/skylake_client (visited on 03/10/2022).

LIST OF FIGURES

- 1 Simplified overview an Intel Skylake CPU [[Gru20](#), fig. 2.1]. For the memory sub-system, detailed knowledge of the load and store buffers, as well as the TLBs, is not required. The same applies to the allocation queue of the frontend. 4

STATEMENT OF AUTHORSHIP

I hereby confirm that the work presented in this bachelor thesis has been performed and interpreted solely by myself except where explicitly identified to the contrary. I declare that I have used no other sources and aids other than those indicated. This work has not been submitted elsewhere in any other form for the fulfilment of any other degree or qualification.

TODO

TODO