UNIVERSITÄT BONN

TODO

TODO

Bachelor Thesis

by

TODO

in fulfillment of requirements for degree
Bachelor of Science (B.Sc.)

submitted to
Rheinische Friedrich-Wilhelms-Universität Bonn
Institut für Informatik IV
Arbeitsgruppe für IT-Sicherheit

in degree course
Computer Science (M.Sc.)

First Supervisor:	TODO	
	University of Bonn	
Second Supervisor:	TODO	
	TODO	
Sponsor:	TODO	
	University of Bonn	

TODO

ABSTRACT

TODO

Contents

1 INTRODUCTION

Contrary to older processors, modern Intel CPUs implement a number of optimization techniques that increase their efficiency. One of which is the concept of out-of-order execution, which takes advantage of the mutual independence of instructions that would normally be executed sequentially. A second optimization technique is called speculative execution and involves the prediction of which branches are taken. With either technique, the CPU might encounter cases where the current CPU state must be rolled back to a previous one. For out-of-order execution, this happens if an instruction raises an exception (illegal memory access, for example). For speculative execution, this happens if a branch is mispredicted. A rollback causes some instructions that are currently being executed (in-flight) to continue execution for a short amount of time.

Even though rollbacks are meant to make sure in-flight instructions do not cause any lasting side effects on the microarchitectural state of the CPU, it was discovered that they can change the contents of caches and other buffers. The disclosure of both Spectre and Meltdown in early 2018 introduced a whole family of vulnerabilities that take advantage of both out-of-order and speculative execution to leak secrets over the processor's caches. And while the performance losses introduced by software and hardware mitigations are measurable, neither vulnerability can be exploited freely on a fully patched system. As a result, however, the process of trying to exploit one of the vulnerabilities for the sake of learning how they work in detail can be challenging. Apart from the software, which may be obtained by installing an older version of an operating system that does not implement any mitigations, one must also make sure their CPU is effected by the vulnerabilities and has not yet received any relevant microcode updates from Intel. Often times, this means that a user's personal computer does not meet these requirements.

We design and implement a graphical CPU emulator that supports single-step out-of-order and speculative execution and is vulnerable to select variants of Spectre and Meltdown. While it is a simplification in comparison to real hardware, the emulator allows its users to gain a better understanding of how exactly the two vulnerabilities work and can be exploited. Furthermore, the user may experiment with (ineffective) mitigations or implement their own microprograms that are executed once rollbacks are completed. We supply example programs that can be run by the emulator and serve both as an entry point for the user as well as the basis of our evaluation.

Firstly, chapter 2 briefly gives an overview of relevant components of vulnerable Intel CPUs, presents the concepts of out-of-order and speculative execution in greater detail, and introduces both Meltdown and Spectre to which the resulting emulator is supposed to be vulnerable. Secondly, chapter 3 further describes the target audience of the emulator and which variants of the vulnerabilities have been chosen, while chapter 4 documents the implementation of the emulator by describing each main component and explaining how rollbacks are implemented. Additionally, it contains an overview of the set of ISA instructions available to the user. Furthermore, chapter 5 explores the

graphical user interface by defining the goals its design is supposed to accomplish and documenting design choices and important features. Chapter 6 provides a demonstration of the emulator, which includes example programs, and determines how effective the implemented mitigations are. Finally, chapter 7 summarizes the other chapters, briefly reflects on how valuable the emulator might be to our target audience, and gives ideas for future improvements.

2 Background

This chapter briefly covers the theoretical background needed to use the presented emulator and continue its development. The reader is assumed to have an understanding of elementary CPU concepts, such as pipelining, caching, and virtual addressing.

2.1 CPU

TODO

2.2 Out-of-order execution

As the name implies, out-of-order execution refers to the idea of executing instructions in a different order than the one in which they are given [Gru20]. With multiple execution units that run in parallel (as described in sec. 2.1), CPUs can take advantage of mutually independent instructions and execute them at the same time.

The basic realization of this concept is provided by Tomasulo's algorithm [Tom67]. It introduces two components, the first of which is the reservation station, which collects the operands of instructions until they are executed by the execution units. Crucially, the corresponding execution units do not need to wait until all operands are present and can instead compute other instructions. The second component of Tomasulo's algorithm is the Common Data Bus (CDB), which connects all reservation stations and execution units. Whenever a result is computed by an execution unit, the result is broadcast onto the CDB and thus made available to all reservation stations that are waiting for it. This important step ensures that results are not written to registers first, just to be read again by other instructions that need them as operands.

Strictly speaking, according to Tomasulo, each set of execution units needs its own reservation station, however, more modern implementations by Intel use a single unified reservation station that handles all types of instructions, rather than just one [Tom67] [Wik].

2.3 Speculative execution

Speculative execution allows a CPU to predict the outcome of comparisons and other branch instructions. This prevents stalls when waiting for the instruction that determines which branch is taken to finish. Similarly to out-of-order execution, rollbacks are needed in some cases. However, in addition to exceptions, they also occur if a branch was mispredicted. [Gru20]

To predict the outcome of branch instructions, CPUs include a branch prediction unit (BPU) [Gru20]. While available in different configurations, many modern CPUs record the most recent outcomes of a branch with a counter [Gru20] that is either incremented if the branch is taken, or decremented.

2.4 MELTDOWN AND SPECTRE

Meltdown [Lip+18] and Spectre [Koc+19] abuse out-of-order and speculative execution to leak data from memory addresses that are normally inaccessible to the attacker over the caches of the CPU.

2.4.1 MELTDOWN

On a high level, Meltdown [Lip+18] works by forcing exceptions when reading data inaccessible to the attacker and transiently encoding this data into the cache to retrieve it once the rollback has completed. What enables Meltdown is a small time window between an invalid memory access and the raising of an exception [Lip+18].

The Meltdown-US-L1 variant of Meltdown [Lip+18] to which the presented emulator is vulnerable works by accessing a memory address for which the attacker has no permission. Firstly, the attacker allocates an oracle array and ensures none of its entries are present in the cache. Upon loading the contents of the inaccessible memory location into a register, the attacker uses the value to access the array at a specific offset. When the rollback caused by the access violation has completed, the attacker measures the access time to each of the array entries to determine which one has been accessed. It is important that the data to be stolen is currently cached. While there are numerous other variants of Meltdown that differ from the basic variant mainly by how they force an exception and from which microarchitectural buffer they leak data, these types of attacks are out of the scope of this work.

2.4.2 SPECTRE

Spectre [Koc+19], on the other hand, relies on the CPU mispredicting a branch and transiently executing instructions that are not part of the correct execution path. This misprediction in a victim process can be induced by maliciously configuring the BPU (sec. 2.3). Depending on the instructions that are wrongfully executed by the victim, traces may later be found in the processor's cache. Similarly to Meltdown, different variants of Spectre exist. The presented emulator is vulnerable to variant 1 of Spectre, which takes advantage of the CPU mispredicting the outcome of comparison instructions.

2.4.3 MITIGATIONS

The mitigations for Meltdown are available in both software and hardware, some of which are implemented in the emulator to allow users to experiment with them and determine their effectiveness. A first, rather simple, mitigation for Meltdown is to disable out-of-order execution [Lip+18], which would completely prevent an attacker from encoding the normally inaccessible data into the cache.

Later revisions of Intel's architectures introduced further mitigations. Although undocumented by Intel, researchers suspect the processor still performs the illegal read, but zeros out the data that is given to dependent instructions before raising an exception [Can+20]. The emulator implements this mitigation, which may be enabled by the user. Other mitigations deployed by operating systems, such as KTPI (or KAISER) [Lip+18], are highly effective, but out of scope.

Unlike Meltdown, Spectre appears to be a design flaw. While one might argue that transiently computing on real values after it is already known an exception will occur is a bug, the behavior abused by Spectre is a direct consequence of speculative execution (sec. 2.3). As a result, an effective yet questionable mitigation is to simply disable speculative execution [Koc+19]. Alternatively, Intel recommends potential victim's of Spectre v1 to use an "lfence" instruction where appropriate, which ensures prior load instructions retire before continuing, thus effectively disabling speculative execution for certain parts of an application [Int18]. A fence instruction is provided as part of the emulator's instruction set.

An additional mitigation that works against both Meltdown and Spectre is to flush the entire cache after a rollback. While highly inefficient, it does prevent the encoding the inaccessible data during the transient execution phase of the attacks. This mitigation can be realized using the emulator by a sequence of instructions that are executed after each rollback.

3 SPECIFICATION OF OUR TASK

todo

4 CPU EMULATOR/ BACKEND

In this chapte, we introduce the backend of our emulator program. It contains the elements of our program that emulate actual CPU components. Our emulator is based on information about modern real life CPUs, especially the x86 Intel Skylake architecture [Wik].

In our program, we model the distinct components of a real life CPU with a modular setup making use of the object oriented functionalities of Python3. Breaking up the source code into individual CPU components also makes it easier to understand and maintain. Additionally, we have made simplifications and modifications in comparison to a real life CPU, so the emulator is as clear and easy to understand as possible while still implementing an actual out of order execution and providing the necessary functionality for Meltdown and Spectre attacks.

In this chapter we firstly introduce the components of our CPU emulator, how they work (together) and which part of a real life CPU they emulate sec. 4.1. Then we explain how our emulator provides out of order execution and how it may differ from the general Tomasulo algorithm introduced in sec. 2.2. Subsequently, we show how we implemented rollbacks and exception handling, especially with regards to how out the implementation allows for Meltdown and Spetre attacks sec. 4.3. Then we give an overview over our instruction set architecture sec. 4.4. Lastly we show how our emulator can be adapted for different demonstrations and attacks without changing the source code via a config file sec. 4.5.

4.1 CPU COMPONENTS AND OUR EQUIVALENTS/ MODELS (10-11 PAGES)

todo

4.1.1 CPU

The primary purpose of the CPU module is to allow all other components to work together. That is, the CPU initializes all other components and provides interfaces to their functions, which allows the GUI to visualize the current state. In addition, the CPU provides functions which load user programs and a tick function that is called each cycle and does the following: instructions from the instruction queue are fetched from the frontend and forwarded to the execution engine, until either no more slots in the unified reservation station are available, or the instruction queue is empty. It then calls the tick function of the execution engine. In case of a rollback, the instruction queue maintained by the frontend is flushed. If the rollback was caused by a faulting load instruction, execution is resumed at the next instruction (as described in sec. 4.3). In case a branch was mispredicted, the frontend is notified and refills the instruction queue accordingly. If configured, corresponding

microprograms are sent to the instruction queue. Lastly, a snapshot of the current state of the CPU is taken.

The second purpose of the CPU is to provide the snapshot functionality which allows a user of the emulator to step back to previous cycles. The snapshot list is simply a list which grows at each cycle, where each entry is a deep copy of the CPU instance. However, this would imply that each snapshot maintains a full list of snapshots, where each entry contains a list of snapshots, and so on. To combat this, there is only a single global list of snapshots, and each snapshot entry maintains a reference and an index to its own entry. As a result, the default deep copy function is adjusted accordingly. Due to the low complexity of the programs we expect users to run, at the moment, no maximum number of available snapshots is configured.

4.1.2 INSTRUCTIONS AND PARSER

todo

4.1.3 DATA REPRESENTATION

todo

4.1.4 CPU FRONTEND

In modern CPUs, the CPU frontend provides an interface between code and execution engine. It contains components to fetch and decode the instructions from a cache and supply them to the CPU in a queue. It is also involved in speculative execution by predicting the result of conditional jumps and supplying instructions to the execution engine accordingly. [Wik]

In our emulator we simplify the components and procedures. We also separate the branch prediction unit (BPU), that manages and predicts the results of conditional jumps, from the rest of the frontend. This makes our code easier to understand and makes potential future changes or additions to the BPU more convenient.

BRANCH PREDICTION UNIT (BPU)

The BPU of modern CPUs plays a vital role in enabling speculative execution. It stores information about previously executed conditional branch instructions and predicts the outcome of future branch instructions accordingly. The rest of the CPU can then speculatively execute further instructions at an address based on the predicted outcome of the branch instead of stalling until the branch instruction is processed by the execution engine. If the prediction was false, a rollback is performed on the speculatively executed instruction sec. 4.3. If the prediction was true, the execution is overall faster than without speculative execution.

Detailed information about the BPU of our modern base CPU is not widely available [Wik]. In general, the Gshare BPU of modern CPU consists of multiple components. It contains a branch target buffer (BTB) that holds the predicted target addresses for conditional branches. It also uses a

pattern history table (PTH), that can be implemented as a 2-bit-saturating counter, and a global history register (GHT) to predict whether a conditional jump will be taken or not. [**SCA**]

In our emulator, we forgo the BTB entirely. Real life CPUs benefit from stored target addresses since they have to expensively decode each branch instruction before they can work with the target address [Wik]. In our emulator, the parser decodes the jump labels from the assembler code and directly provides them to the frontend, so storing them in an additional buffer is unnecessary sec. 4.1.2. We further forgo the GHT because it is not strictly necessary to execute a Meltdown or Spectre attack. Additionally, our emulator and its behaviour are easier to understand and predict without it, which is important when implementing microarchitectural attacks for didactic purposes. The BPU of our emulator only consists of a PHT, which is enough for simple Spectre-PHT variants [**reference-eval-spectre**].

The default PHT used in our emulator holds an array called counter of configurable length to store several predictions. The instructions are assigned to different prediction slots by the last bits of their index in the instruction list. For each of the slots, the prediction can take the four different values from zero to three, where zero and one indicate that the branch will probably not be taken and two and three indicate that the branch is likely to be taken.

The source code for our emulator also contains a more simple BPU with only one slot for all instructions. Since the number of slots in the default BPU is freely configurable by the user, this simple BPU is now obsolete.

When the BPU is updated with an actual branch outcome from the execution engine, the prediction in the PHT is updated by a 2-bit counter. This means, that if the prediction was right, the counter remains at or updates to zero (strongly not taken) or three (strongly taken) respectively. If the prediction counter is at zero but the branch is actually taken, the counter is updated to one (weakly not taken). If it is at a one when the branch is taken, it is directly updated to three. The counter behaves similarly when it has tha value two or three and the branch is not actually taken.

Instruction Queue

In a real life CPU, the overall purpose of the frontend is to provide the execution unit with a steady stran of instructions so the backend is busy as much as possible and therefore efficient. In a modern x86 CPu, the frontend has to fetch x86 macro-instructions from a cache and decode, optimize and queue them repeatedly to provide the backend with a queue of μ-instructions ready for issuing in the execution engine. [Wik]

In our emulator, except for the BPU, the functionality of the CPU frontend is bundled in frontend.py. It is significantly simplified compared to a real life x86 CPU, especially since the we use only one type of instructions instead of distinguishing between macro- and μ-instructions sec. 4.4. They are already provided as a list by the parser, which renders the decoding and optimization steps in the frontend unnecessary sec. 4.1.2.

The main task of our frontend is to acts as interface between instruction list provided by the parser and the execution engine sec. 4.1.6. It provides and manages the instruction queue, which holds the instructions that the execution engine should issue next. Conditional branches with their respective

BPU predictions are taken into account when filling the queue. This enables speculative execution which is needed for Spectre attacks sec. 2.4.

The central component of our emulated frontend ist the instruction queue. In our version, it does not only hold the instructions themselves, but also for every instruction in the queue, the respective index in the instruction list is stored. For branch instructions it also holds the respective branch prediction from the BPU at the time that the instruction was added to the queue. This additional information is needed by the execution engine to handle mispredictions and other exceptions sec. 4.1.6.

When adding instructions to the queue, the frontend selects them from the instruction list, adds the additional information for the execution engine and places them into the instruction queue until the queue's maximum capacity is reached. The frontend maintains a program counter (pc) that points to the next instruction in the list that should be added to the queue. When the frontend encounters a branch instruction and the branch is predicted to be taken, the frontend adjusts the pc to resume adding instructions at the branch target. If a branch was mispredicted, the frontend provides a special function to reset the pc and refill the instruction queue with the correct instructions.

Additionally, the frontend provides a function to add a μ-program to the queue. It consists of a list of instructions separate from the parser instruction list. When adding the μ-program to the queue, the frontend may exceed the maximum queue capacity. This functionality can be used to implement mitigations against microarchitectural attacks, e.g. by adding a μ-program as part of the exception handling after an illegal load [**reference-eval-mitigations**].

The frontend provides interfaces to both read and take instructions from the queue. It also provides a function that combines taking an instruction from the queue and refilling it, . Additionally, the frontend has an interface for flushing the whole queue at once without taking the instructions from the queue. This can be used when demonstrating mitigations against microarchitectural attacks [**reference-eval-mitigations**].

The frontend provides further basic interfaces, e.g. for reading the size the instruction queue and reading and setting the pc. These are used by the other components during regular execution, e.g. when issuing instructions to the execution engine, but also to reset the queue to a certain point in the program after an exception has occured sec. 4.3. Since our emulator only executes one program at a time, the other components can check via another interface whether the frontend has reached the end of the program.

4.1.5 MEMORY

todo

4.1.6 EXECUTION ENGINE

todo

4.2 Out of Order Execution (2 pages)

todo

4.3 Exceptions and Rollbacks (2-3 pages)

todo

4.4 ISA (2 pages)

Real life Intel x86 CPUs differenciate between two types of instructions or operations. Macro-operations refer to the relatively easily human readable and convenient but complex instructions that are described by the x86 ISA. Their length differs between the instructions. Internally, in the execution units, the CPU works on μ-operations, which are small operations of a fixed length. One macro-operation contains one or multiple μ-operations. The CPU frontend has to decode the macro-operations into μ-operations in an expensive multi step process. sources: [Wik], https://en.wikichip.org/wiki/macro-operation, https://en.wikichip.org/wiki/micro-operation

Our CPU emulator only uses one type of instructions. They are directly read from our assembler code by the parser and passed to the execution engine without further decoding, splitting or replacing sec. 4.1.2, sec. 4.1.4. To show basic Meltdown and Spectre variants, we do not need overly complex instructions, e.g. instructions that contain multiple memory accesses in one or that are used to perform encryption in hardware [ref_evaluation_meltdown], [ref_evaluation_spectre]. Basic arithmetic operations, memory accesses, branches and a few special operations are sufficient for the demonstrated attacks and are both easy to implement as single instructions and to use in assembler code that should be well understood by the author. Using the same operations throughout the emulator also makes the visualization more clear and easier to follow, e.g. when the same operations appear, one after the other, in the visualization of the assembler code, the instruction queue and the reservation stations [ref_ui].

4.4.1 Default Instruction Set

In order that our CPU emulator can recognize and work with an instruction, it has to be registered with the parser sec. 4.1.2. In our default setting default, we register a basic set of instructions with the parser so students can start writing assembler code and using the emulator right away. This basic instruction set is also used in our example programs in [ref_UI].

Our relatively small instruction set is based on a subset of the RISC-V ISA. It offers a selection of instructions that is sufficient to implement Meltdown and Spectre attacks as well as other small assembler programs while still being of a manageable size so students can start to write assembler code quickly without spending much time to get to know our ISA. The syntax of the assembler representation is also based on RISC-V (as introduced in the "RISC-V Assembly Programmer's Handbook" chapter of the RISC-V ISA) [ref_RISC-V]. If needed, students can add further instructions by registering them with the parser sec. 4.1.2.

In the following subchapters we introduce the instructions of our default ISA. They are grouped according to their respective instruction type in the emulator except for the special instructions which are grouped together sec. 4.1.2.

Arithmetic and Logical Instructions without Immediate

These are basic arithmetic and logical instructions that operate solely on register values, i.e. both source operands and the destination operand reference registers. For simplicity, we write, for example, Reg1 when referring to the value read from or stored in the register referenced by the first register operand.

Each of these default operations uses the respective python standard operator on our Word class to compute the result, except for the right shifts. For the logical and the arithmetic right shift, the python standard right shift operator is used on the unsigned and the signed version of the register value respectively. When returning the result as a Word, it is truncated to the maximal word length by a modulo operation, if necessary sec. 4.1.3. This means, that any potential carry bits or overflows are effectively ignored.

Arithmetic and Logical Instructions without Immediate		
Instr. Name	Operators	Description
add	Reg1, Reg2, Reg3	Reg1 := Reg2 + Reg3
sub	Reg1, Reg2, Reg3	Reg1 := Reg2 − Reg3
sll	Reg1, Reg2, Reg3	Reg1 := Reg2 << Reg3
srl	Reg1, Reg2, Reg3	Reg1 := Reg2 >> Reg3 logical
sra	Reg1, Reg2, Reg3	Reg1 := Reg2 >> Reg3 arithmetical
xor	Reg1, Reg2, Reg3	Reg1 := Reg2 xor Reg3
or	Reg1, Reg2, Reg3	Reg1 := Reg2 or Reg3
and	Reg1, Reg2, Reg3	Reg1 := Reg2 and Reg3

Arithmetic and Logical Instructions with Immediate

These are basically the same instructions as in sec. 4.4.1. The main difference is, that the second source register is replaced by an immediate operand which is set directly in the Assembler code. This immediate is used as the value of a Word in the execution engine, so it is truncated by a modulo operation to be in the appropriate range sec. 4.1.3, sec. 4.1.6.

Arithmetic and Logical Instructions with Immediate		
Instr. Name	Operators	Description
addi	Reg1, Reg2, Imm	Reg1 := Reg2 + Imm
subi	Reg1, Reg2, Imm	Reg1 := Reg2 − Imm
slli	Reg1, Reg2, Imm	Reg1 := Reg2 << Imm
srli	Reg1, Reg2, Imm	Reg1 := Reg2 >> Imm logical
srai	Reg1, Reg2, Imm	Reg1 := Reg2 >> Imm arithmetical
xori	Reg1, Reg2, Imm	Reg1 := Reg2 xor Imm
ori	Reg1, Reg2, Imm	Reg1 := Reg2 or Imm
andi	Reg1, Reg2, Imm	Reg1 := Reg2 and Imm

Memory Instructions

Memory Instructions basic memory interactions load and stores both word-wise for convenience and byte-wise for the fine granular (?wording?) access needed in micro architectural attacks flush flushes (sets data and tag (if no data at tag left) to none) a cacheline for more detailed information about how these affect the memory and especially the cache, look at the backend.memory subchapter address calculation always the same: addr = Reg + immediate one acts as base, the other as offset why this way around? fixed base and variable offset seem more logical ask Jan-Niklas, maybe different logic or real life model

Branch Instructions

Branch Instructions same basic structure: comparison on 2 registers if evaluates to true: resume execution at a predefined label in the program if fales: next instruction labels are automatically resolved by the parser (reference parser) multiple options for the condition, so students can choose what suits them best again address calculation always the same:

Special Instructions

Special Instructions rdtsc important if one wants to try CBSCA without just looking at the memory visualization basic timing instruction cyclecount: number of executed execution unit ticks (ref. execution unit) all the ticks or reset when the instruction is called? -> look into code fence can be used for a sort of mitigation compare to mfence, lfence etc. in x86 all instr in the EU unit at the point of issueing the fence are ex- ecuted before the fence ist executed; no new instructions are issued before the fence is executed

ggf. auf Cheat Sheet im Anhang verweisen maybe put information like address calculation in table description so everyone has all the information adjust cheat sheet and table snippets so the wording is nice and the table is as non-redundant as possible

Instructions			
Instr. Kind	Instr. Name	Operators	Description
ALU Instructions without Immediate	add	Reg1, Reg2, Reg3	Reg1 := Reg2 + Reg3
	sub	Reg1, Reg2, Reg3	Reg1 := Reg2 − Reg3
	sll	Reg1, Reg2, Reg3	Reg1 := Reg2 << Reg3
	srl	Reg1, Reg2, Reg3	Reg1 := Reg2 >> Reg3 logical (fill space with 0)
	sra	Reg1, Reg2, Reg3	Reg1 := Reg2 >> Reg3 arithmetical (fill space with sign bit)
	xor	Reg1, Reg2, Reg3	Reg1 := Reg2 xor Reg3
	or	Reg1, Reg2, Reg3	Reg1 := Reg2 or Reg3
	and	Reg1, Reg2, Reg3	Reg1 := Reg2 and Reg3
ALU Instructions with Immediate	addi	Reg1, Reg2, Imm	Reg1 := Reg2 + Imm
	subi	Reg1, Reg2, Imm	Reg1 := Reg2 − Imm
	slli	Reg1, Reg2, Imm	Reg1 := Reg2 << Imm
	srli	Reg1, Reg2, Imm	Reg1 := Reg2 >> Imm logical (fill space with 0)
	srai	Reg1, Reg2, Imm	Reg1 := Reg2 >> Imm arithmetical (fill space with sign bit)
	xori	Reg1, Reg2, Imm	Reg1 := Reg2 xor Imm
	ori	Reg1, Reg2, Imm	Reg1 := Reg2 or Imm
	andi	Reg1, Reg2, Imm	Reg1 := Reg2 and Imm
Memory Instructions	lw	Reg1, Reg2, Imm	addr:=Reg2(base)+Imm(offset), Reg1:=(Mem_word[addr])
	lb	Reg1, Reg2, Imm	addr:=Reg2(base)+Imm(offset), Reg1:=(Mem_byte[addr])
	sw	Reg1, Reg2, Imm	addr:=Reg2(base)+Imm(offset), Mem_word[addr]:=Reg1
	sb	Reg1, Reg2, Imm	addr:=Reg2(base)+Imm(offset), Mem_byte[addr]:=Reg1
	flush	Reg, Imm	addr:=Reg(base)+Imm(offset), flush_cashline(adr)
Branch Instructions	beq	Reg1, Reg2, Label	pc:=Label if Reg1==Reg2, else pc+ =1
	bne	Reg1, Reg2, Label	pc:=Label if Reg1! =Reg2
	bltu	Reg1, Reg2, Label	pc:=Label if u(Reg1)<u(Reg2)
	bleu	Reg1, Reg2, Label	pc:=Label if u(Reg1)<=u(Reg2)
	bgtu	Reg1, Reg2, Label	pc:=Label if u(Reg1)>u(Reg2)
	bgeu	Reg1, Reg2, Label	pc:=Label if u(Reg1)>=u(Reg2)
	blts	Reg1, Reg2, Label	pc:=Label if s(Reg1)<s(Reg2)
	bles	Reg1, Reg2, Label	pc:=Label if s(Reg1)<=s(Reg2)
	bgts	Reg1, Reg2, Label	pc:=Label if s(Reg1)>s(Reg2)
	bges	Reg1, Reg2, Label	pc:=Label if s(Reg1)>=s(Reg2)
Special Instructions	rdtsc	Reg	Reg:=cyclecount(number of executed execution unit ticks)
	fence	none	all instr in the EU unit at the point of issueing the fence are executed before the fence ist executed; no new instructions are issued before the fence is executed

4.5 Config Files (1 page)

todo

5 User Interface and Usage

todo

6 EVALUATION

todo

7 CONCLUSION

todo

8 Catchy Lab title

8.1 Abstract

8.2 Introduction (1 page)

```
1  general task/ goal
2      CPU emulator with actual out of order execution, maybe speculative
       execution and comprehensible implementation and documentation
3      implement different µ-architectural features like out-of-order exe.,
       speculative exe. b.c. we want to observe and teach µ-arch. attacks
4  brief context of the task/ goal (e.g. a sentence on Meltdown and why it is
       important to understand it)
5  structure of the report
```

8.3 Brief Theoretical Background (2-3 pages)

```
1  premise according to Felix: reader knows SCA lecture -> brief recaps/
       reminders to reference back to from the other chapters
2
3  ### really brief introduction to CPU
4
5      multiple components
6          maybe mostly via picture
7
8      how they work together
9
10     maybe data-flow from instruction to result
11
12     a lot of hypothticals due to trade secrets
13
14     suggested literature:
15         maybe Gruss Diss. and other works on cbsca
16         maybe textbooks
17         maybe CPU wiki for pictures
18
19 ### out-of-order execution
20
21     a bit more in depth b.c. this is not content of the sca lecture
22
23     maybe why do we need it/ advantages
```

```
24
25      brief explanation of Tomasulo and how it works
26
27      suggested literature:
28          original Tomasulo paper from 1965
29          maybe a more recent/ didactically edited explanation (is this in
        texbooks?)
30
31      also something about speculative execution
32
33  ### Meltdown and Spectre
34
35      brief introduction to Meltdown (and Spectre) in general
36
37      slightly more in-depth introduction to the Meltdown attack we picked
        and want to run in our emulator
38
39      highlight the relevant parts of the CPU, maybe mention how they
        interact in Meltdown (and Spectre)
40
41      mitigations and maybe how successfull they are
42
43      suggested literature:
44          ggf. Dissertation Gruss if we mention basic cbsca
45          Gruss et al. papers on Meltdown and Spectre e.g. https://gruss.cc/
        files/meltdown.pdf or https://gruss.cc/files/meltdown_cacm.pdf
46          Canella, Gruss et al. on mititgations https://gruss.cc/files/
        transient-attacks.pdf
```

8.4 maybe Specification of the task (1 page)

```
1  emulator to execute and teach Meltdown
2      who is the target audience
3      which Meltdown attacks specifically
4      etc. further concretisations
5
6  are there existing solutions/ related works?
7      Felix' old emulator? citable?
8      suggested literature:
9          todo
```

8.5 CPU emulator/ backend (17-18 pages)

```
1  maybe general info e.g. that we wrote it in python, if we used special
        tools/ libraries
2  for everything we implemented:
```

```
 3        which part of a real life cpu does this model/ how is this done in
     real life cpus
 4        briefly how does it work
 5        why did we choose to model it like we do
 6            nice code structure/ code easy to understand by students
 7            some features more or less relevant for meltdown
 8            etc.
 9        maybe what did we leave out
10        challenges?
11
12 ### CPU Components and our equivalents/ models (10-11 pages)
13
14    modular setup based on real life CPUs and nice coding conventions
15
16    #### CPU
17
18        contains/ controls the rest
19
20        preparation (loading the program, init the rest)
21
22        ticks
23
24        coordinate rollbacks?
25
26    #### Instructions and Parser
27
28        instructions
29
30        parser
31
32    #### Data representation
33
34        how we model and handle data
35
36        byte, word
37
38    #### CPU frontend
39
40        Branch Prediction Unit (BPU)
41
42        Instruction Queue
43
44    #### memory
45
46        no virtual adresses
47
48        mmu
49
50            how do we store data
```

```
51
52            how do we model data access (i.p. wrt Meltdown)
53
54        cache
55
56            what kinds of caches can we represent
57
58            how do we model if something is cached and the access times (i.
      p. wrt Meltdown)
59
60      #### Execution Engine
61
62          Reservation Station/ Slots
63
64          unlimited execution units
65
66  ### Out of Order Execution (2 pages)
67
68      our version of out-of-order execution/ Tomasulo
69
70      where in our program do we implement which components
71
72      what did we leave out/ do differently
73
74  ### Exceptions and Rollbacks (2-3 pages)
75
76          in particular wrt making Meltdown possible
77
78          general goal/ concept of rolling back after a misprediction or an
      exception
79
80          snapshots and interaction of the CPU components
81
82  ### ISA (2 pages)
83
84      overview of ISA
85
86      μ-instr. vs. ISA
87          mixture of both in one: only μ-code, but more abstract than in real
       life
88          do not need to think about page faults etc. anyways, do not need
      overly complex instructions
89
90      reasoning behind the choice of instructions (manageable size and
      instructions (e.g. no divide by zero) balanced with functionality
      particularly wrt Meltdown)
91
92
93  ### config files (1 page)
```

```
94
95       what can be configured without changing the source code
96           why these variables? relevant for Meltdown?
```

8.6 Our Visualisation and Usage/ Frontend (10 pages)

```
1  ggf. gemäß Anmerkung von Lenni umsortieren: die idee das UI nicht mit in
       die 'unsere designetnscheidungen' zu nehmen, sondern im prinzip so als
       Manual abzukapseln finde ich eigentlich ganz gut. Müssen wir aber dann
       mal in der Praxis schauen. Soll ja keine didaktischen begründungen das
       manual zu sehr aufblähen, vllt wird das sinnvoll, dann einen teil des
       UI im "backend" kapitel einzubauen, und dann wirklich ein cleanes
       manual kapitel zu haben
2
3  ### general concept
4
5      goal: UI for the emulator with visualisation of CPU/ memory components
       and their contents
6
7      in terminal
8          maybe comparison to existing analysis tool/ debugger gdb
9
10     triggers/ controls the actual emulator
11         overview features, e.g. breakpoints, step-by-step and stepback
12
13     maybe which tools/ libraries were used?
14
15 ### features in more detail and their didactic purpose
16
17     breakpoints, step-by-step, stepback and more
18
19     challenges/ design choices during the implementation
```

8.7 Demonstration/ evaluation (7 pages)

```
1  what kind of system do we need/ did we use to run this?
2      which python Version?
3
4  ## general demonstration
5
6      brief example program showing all the features in a "normal" execution,
       e.g. adding stuff
7
8  ### Meltdown und Spectre demonstration
9
10     #### Example Program Meltdown
```

```
11
12          show example program
13
14          maybe compare to example program for real life architecture from
       SCA or literature (Gruss) if available
15
16          explain which Meltdown variant it implements
17
18          briefly highlight which components (we expect to) interact to make
       it work
19
20          how well does it work?
21
22      #### maybe example program spectre
23
24          same as Meltdown
25
26      #### mitigations
27
28          what is possible in our program as is
29              planned:    cache flush: microcode -> config file
30                          mfence im assembler (normally in compiler)
31                          aslr directly in program -> config
32                          flush IQ ->
33                          disable speculation und out of order (nice to have)
       -> config
34
35          how effective are these
36              in real life
37              in our program
38
39          what would be the necessary steps/ changes to the program for
       further mitigations
40              compare to changes in hardware by the manufacturers
```

8.8 Conclusion (1 Page)

```
1  recap goals
2
3  did we reach the goal of implementing a CPU emulator where a user can
       perform a Meltdown attack
4
5  how many Meltdown attacks are possible?
6
7  is Spectre possible?
8
9  mitigations
10      which mitigations did we implement
```

```
11      is this a good amount/ sample of real world mitigations or do we miss
        important ones?
12      how well do they perform (also in comparison to how they perform in the
         real world)
13
14  value to students:
15      how easy to use and convenient do we think our program is?
16      do we think this will be a good tool for teaching?
17      how accessible is it wrt different host architectures?
18
19  further work
20      more (detailed/ realistic) functionality for more Meltdown and Spectre
        variants
21          more elaborate BPU with btb (and ghr) for more spectre variants
22      more mitigations
23      maybe nice to haves wrt to the visualisation/ general UI functionality/
         ISA?
```

REFERENCES

[Can+20]   CANELLA, Claudio et al.: "KASLR: Break it, fix it, repeat". In: *Proceedings of the 15th ACM Asia Conference on Computer and Communications Security*. 2020, pp. 481–493.

[Gru20]   GRUSS, Daniel: "Transient-Execution Attacks". 2020. URL: https://gruss.cc/files/habil.pdf (visited on 03/11/2022).

[Int18]   INTEL: *Intel Analysis of Speculative Execution Side Channels*. 2018. URL: https://newsroom.intel.com/wp-content/uploads/sites/11/2018/01/Intel-Analysis-of-Speculative-Execution-Side-Channels.pdf (visited on 03/11/2022).

[Koc+19]   KOCHER, Paul et al.: "Spectre attacks: Exploiting speculative execution". In: *2019 IEEE Symposium on Security and Privacy (SP)*. 2019.

[Lip+18]   LIPP, Moritz et al.: "Meltdown: Reading Kernel Memory from User Space". In: *27th USENIX Security Symposium (USENIX Security 18)*. 2018.

[Tom67]   TOMASULO, Robert M: "An efficient algorithm for exploiting multiple arithmetic units". In: *IBM Journal of research and Development* 11.1 (1967), pp. 25–33.

[Wik]   WIKICHIP: *Skylake Microarchitecture*. URL: https://en.wikichip.org/wiki/intel/microarchitectures/skylake_(client) (visited on 03/10/2022).

Statement of Authorship

I hereby confirm that the work presented in this bachelor thesis has been performed and interpreted solely by myself except where explicitly identified to the contrary. I declare that I have used no other sources and aids other than those indicated. This work has not been submitted elsewhere in any other form for the fulfilment of any other degree or qualification.

TODO

_____
TODO